

Rate Limiter

Thursday, February 19, 2026 6:22 PM

Rate Limiter

Brief - It handles the req from client. It just like the traffic controller for the backend infra
- it allows the defined number of requests to the infra for the specified interval of time

Bouncer Metaphor

- club has the capacity (backend infra)
- People are waiting in queue(request)
- Bouncer - only 5 people can enter to the club in every 2 min(job of rate limiter)

Benefits

- it protects the infra from dDos attacks

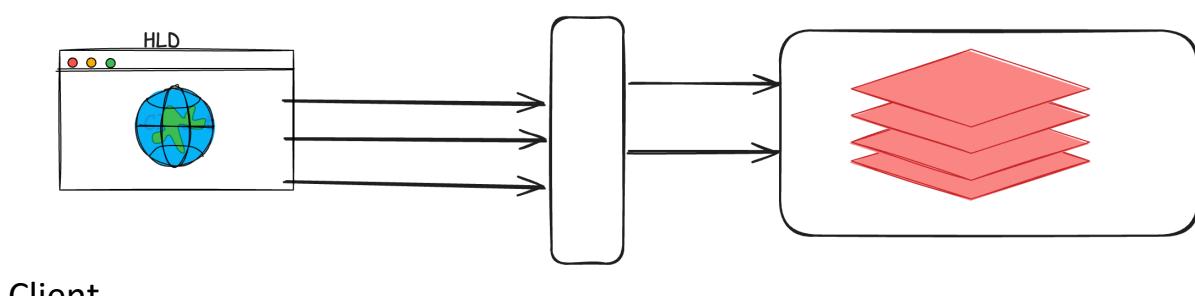
Functional Requirements

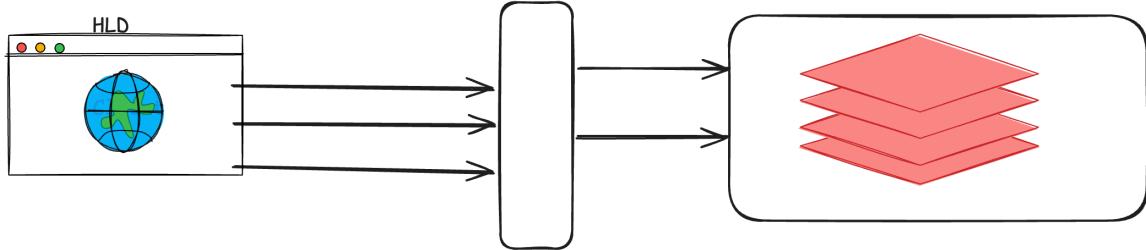
- Limit Req - drop req, if exceeds threshold
- id - Rate limit based on ip, userid or api key
- Feedback - Return a clear error e.g 429 - too many req

NFR

- Low Latency - it should not slow down the actual api call. Either Yes or NO in 10^{-6} sec
- High Availability - if it is down, it should not crash the backend
- Distributed Accuracy - the limit should be global. e.g in a cluster of 10 servers, 100 req/sec should be global not per server.
- Memory efficient - It should store as less data as possible (e.g counter and timestamp)

HLD





Client

The Core Algo

it can use any one from the below based on the business requirement

1. Token Bucket -

- Bucket holds N token. Each req take one. Bucket will be filled at const rate
- e.g. if a user can make 2 posts/day , 20 comments/day , add 150 frnds/day , here user needs 3 different buckets based on the maxtoken size for each bucket

- e.g if we are allowing 100 req for ip , then each ip address needs a bucket
- Pro - Allows burst of traffic
- con - two variables are in the system
 - bucket size
 - refill rate

So it could be tough to tune

2 . Leaking Bucket -

- Same as Token buket , req enters the bucket and processed at fixed rate
- FIFO queue. If q is full, req drops else added

- Pro - Memory efficient , as q size is limited
 - due t fixed rate processing, it is needed where fixed rate is req
- Con - Bursty traffic drops immidiately
 - two params - bucket size and q , so tune could not be easy

3. Fixed Window -

- Counts req is fixed timeframe
- pro - easy to implement
- con - for the spike, it could allow more than the allowed number

4. Sliding Window log -

- It fixes the issue of fixed window

pro - good accuracy
con - consume memory , as timestamps are being logged in the Redis even for the rejected req

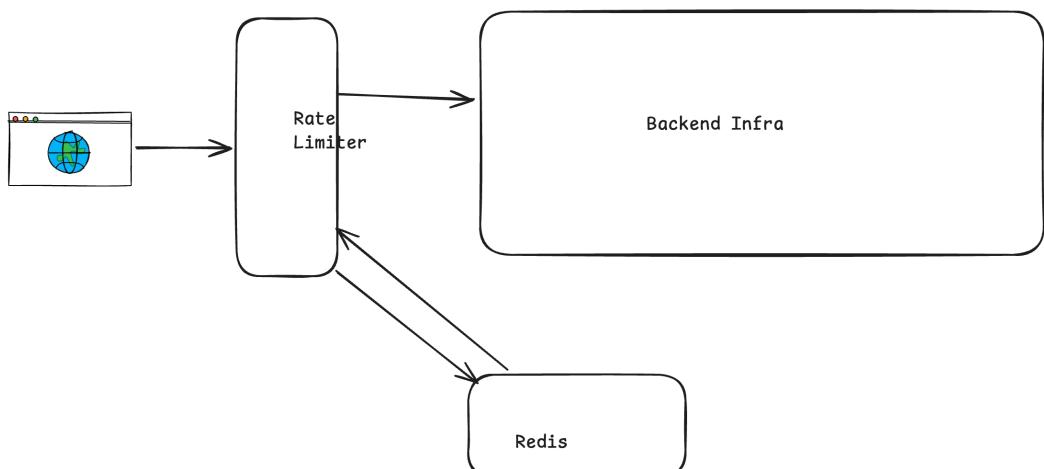
5. Sliding Window counter -

- its a hybrid fixed window+Weighted avg

$$\text{req \#} = \text{curr} + (\% \text{ of the prev window} * \text{total no in the prev window})$$

pro - very efficient. handles traffic spike

Con - as it is an approximation of the prev window, so wont work efficiently if the req is strict look back



- 429 - if a req is rate limited , then the http code it returns

Challenges in Distributed Env

- Race Condition -

- if both req acts on the Redis to get the counter and increment the val by one, so it will be counter+1 however the

correct val would be counter+2.

- Solution - Lock - But it will slow down the system.

- better soln - Lua Script

- sorted set DS in Redis

- Synchronization issue

- in distributed arch , multiple clients and multiple rate limiter. As the web tier is Stateless, if a client will

send request to different rate limitier could cause problem

- Solution - Sticky session - but it is not scalable nor flexible.

Rate limitier has centralized data store like Redis

- Best soin - use a centralized data store like Redis

Consider Latency

- it is best to set up multi DC Rate limiter. Edge location is preferable

Observability -

- Check Rate limiter algo is efficient
- if it failed to serve the traffic spike which are business critical , then need to reconsider the rules

Some suggestion on the client side

- use client cache to avoid being rate limited
- understand the limit and be polite on sending req

